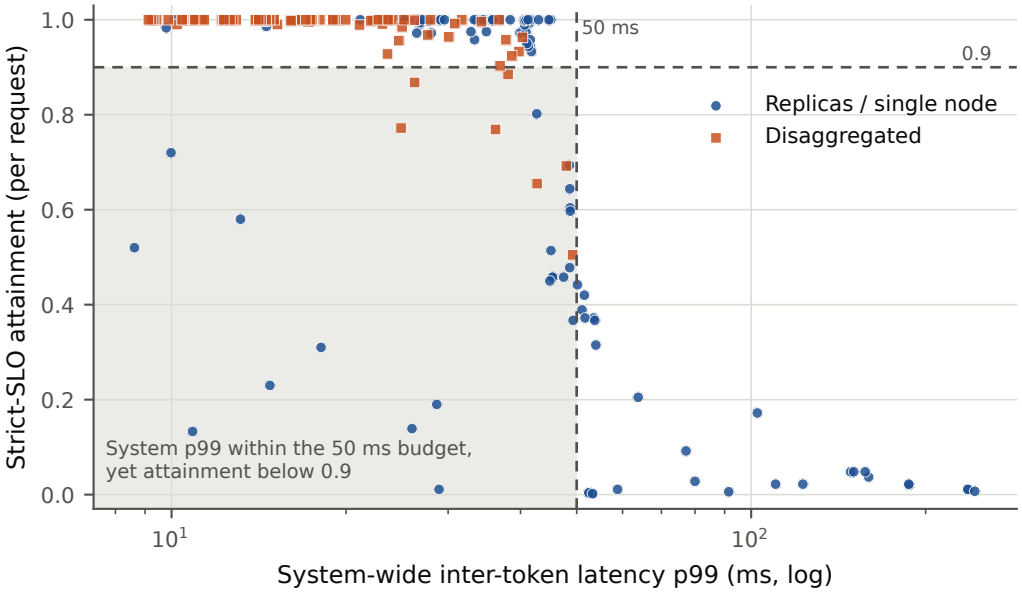
System percentiles hide per-request stalls
1.0
0.8
0.6
0.4
0.2
0.0
Strict-SLO attainment (per request)
50 ms
0.9
Replicas / single node
Disaggregated
System p99 within the 50 ms budget,
yet attainment below 0.9
10^1
10^2
System-wide inter-token latency p99 (ms, log)