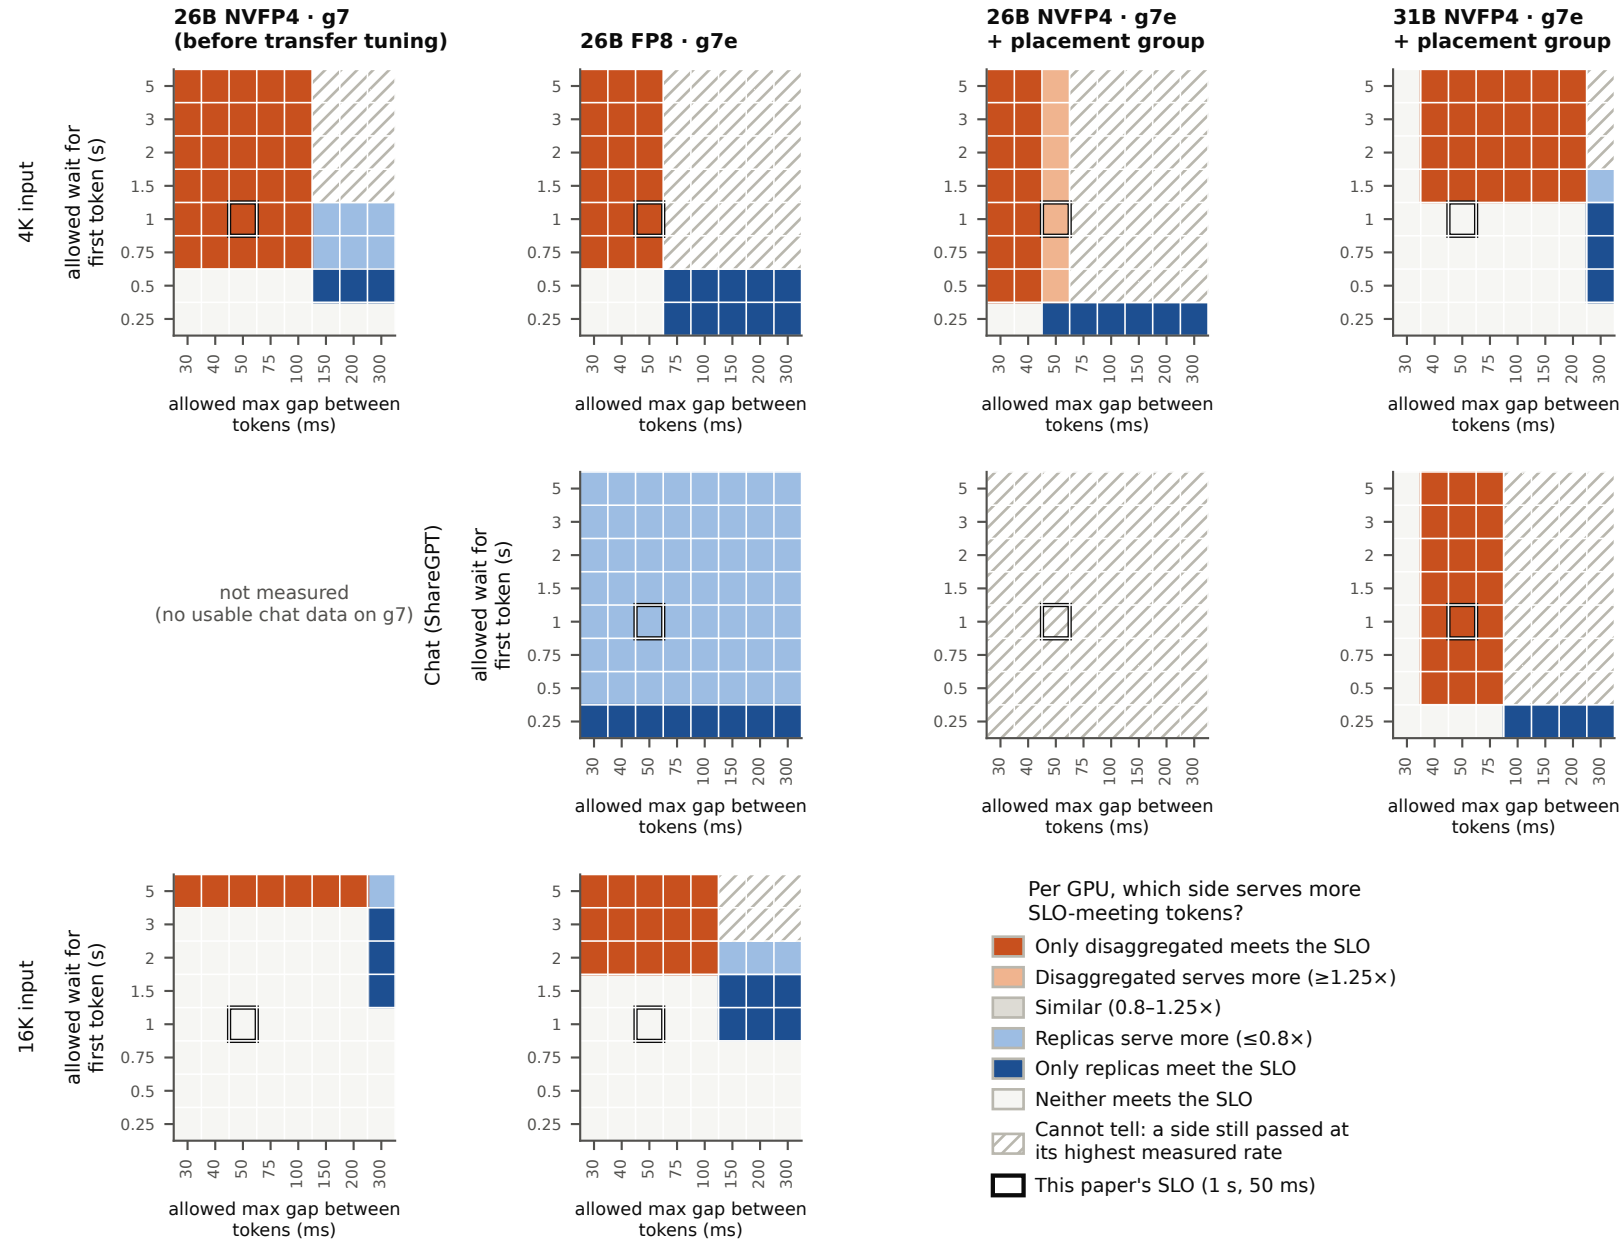
Who wins as the SLO changes: each cell is one SLO (first-token limit × max token-gap limit)
26B NVFP4 · g7
(before transfer tuning)
26B FP8 · g7e
26B NVFP4 · g7e
+ placement group
31B NVFP4 · g7e
+ placement group
4K input
Chat (ShareGPT)
16K input
allowed wait for first token (s)
allowed max gap between tokens (ms)
not measured
(no usable chat data on g7)
Per GPU, which side serves more SLO-meeting tokens?
Only disaggregated meets the SLO
Disaggregated serves more (≥1.25×)
Similar (0.8–1.25×)
Replicas serve more (≤0.8×)
Only replicas meet the SLO
Neither meets the SLO
Cannot tell: a side still passed at its highest measured rate
This paper's SLO (1 s, 50 ms)